

The Agentic Historian — Technical Design Document

Generated on 2025-12-05

The Agentic Historian

An Agentic, Evidence-Grounded Historical Fact Verification Chat Application (LangGraph + Google Search API + Fine-Tuned LoRA LLM)

Course: Deep Learning & Applied LLMs

Due: 8 days

Repo Codename: agentic-historian

1. Project Overview

The Agentic Historian is an AI-powered chat application that helps users **verify historical facts** by combining:

1. **Live evidence retrieval** via the **Google Custom Search JSON API** (no scraping; we only use the official API responses).
2. An **agentic workflow** orchestrated with **LangGraph** to enforce a structured, auditable reasoning pipeline.
3. A **custom fine-tuned LLM** (Mistral/Llama family) using **PEFT/LoRA** trained on **TruthfulQA**-style supervision to improve truthfulness/refusal behavior (and reduce unsupported claims).
4. A **Streamlit** UI that presents: the final answer, the evidence used, source citations, and model uncertainty.

Core product promise:

Answers must be **evidence-grounded**, **citation-backed**, and **explicit about uncertainty** when evidence is insufficient.

2. System Architecture

2.1 High-Level Flow

We implement a **deterministic agent pipeline** using LangGraph:

Router → Researcher (Google Tool) → Fact Analyst → Writer (Fine-Tuned Model)

- **Router:** classifies the user input and selects the appropriate path (historical fact-check vs. out-of-domain vs. ambiguous).
- **Researcher:** issues a Google Search API call and returns structured evidence candidates.
- **Fact Analyst:** evaluates source credibility signals, extracts candidate claims, detects contradictions, and builds a compact, citation-indexed “Evidence Bundle”.

- **Writer:** a fine-tuned LoRA model generates the final answer strictly from the Evidence Bundle and must cite sources.

2.2 Architecture Diagram (Logical)



2.3 LangGraph as the Orchestration Backbone

We use LangGraph because it models agent workflows as graphs with explicit **state**, **nodes**, and **edges**, supporting durable execution and structured routing.

Critically, LangGraph state can be defined as a **Pydantic model**, making the workflow type-safe and debuggable end-to-end.

3. Detailed Component Design

3.1 Shared State (Pydantic)

We keep the entire run auditable using a typed Pydantic state, e.g.:

- `user_query: str`
- `route: Literal["fact_check", "clarify", "out_of_scope"]`
- `search_queries: list[SearchQuery]`
- `search_results: list[SearchResult]`
- `evidence_bundle: EvidenceBundle`
- `final_answer: str`

- citations: list[Citation]
- confidence: float
- run_metadata: dict (latency, token usage, API quota info)

Design rule: every node must be a pure transformation over state:

input state → output state delta

3.2 Router Node

Goal: decide if we run the full fact-check pipeline.

Heuristics (fast + reliable):

- If query is not historical (“write me a poem”), route to `out_of_scope`.
- If query is underspecified (“Did it happen?” with no subject), route to `clarify`.
- Else route to `fact_check`.

Optional: use a lightweight classifier prompt (small base model) but we keep it deterministic when possible.

3.3 Researcher Agent (Google Tool Node)

Tool: Google Custom Search JSON API (Programmable Search Engine).

Constraints: No scraping. We only use:

- title
- snippet
- display link / URL
- optional metadata returned by API

Query strategy:

- Expand user query into 1–3 tightly scoped search queries:
- include dates, names, and disambiguators
- optionally restrict to trusted domains (`site: .edu`, `site: .gov`, museums, encyclopedias) depending on claim type

Output: `SearchResult[]` normalized to a strict schema.

Caching: store results keyed by `(query, date_bucket)` to reduce cost and improve demo reliability.

3.4 Fact Analyst Node (Verification & Evidence Synthesis)

Goal: convert noisy search snippets into a structured evidence pack suitable for a constrained writer.

Steps:

1. **Source scoring** (heuristic):

- Domain type signals (e.g., .gov, .edu, major encyclopedias, museums)
- Result consistency across multiple sources
- Recency relevance (for “as of” queries)

2. Claim extraction:

- Identify atomic claims (date, person, event, location).

3. Contradiction detection:

- If multiple high-ranked sources disagree, mark as “contested”.

4. Evidence bundle construction:

- Build a compact, numbered set of evidence items:
- E1: [source title] – snippet – URL
- E2: ...

Output: EvidenceBundle containing:

- evidence_items: list[EvidenceItem]
- findings: list[Finding] (each finding references evidence IDs)
- verdict: Literal["supported", "not_supported", "contested", "insufficient"]

3.5 Writer Node (Fine-Tuned LoRA Model)

Goal: produce a final response **only from the Evidence Bundle**.

Writer constraints:

- Must cite evidence by ID ([E1], [E2]).
- Must explicitly state uncertainty:
- “Evidence is insufficient to confirm...”
- “Sources disagree...”
- Must refuse to invent missing facts.

Prompt design: a strict instruction template + evidence pack + formatting policy:

- Output format:
- **Answer**
- **Confidence** (0–1)
- **Evidence** (IDs used)
- **Notes / Limitations**

Model: base LLM (Mistral/Llama) + LoRA adapter trained with TruthfulQA-style supervision to reinforce truthful/refusal behavior.

4. Tech Stack Rationale

4.1 uv for Project & Dependency Management

We use `uv` because it provides:

- fast dependency resolution and installation
- a **project workflow** with `uv init`, `uv add`, and automatic lock/sync mechanics
- a universal lockfile and reproducible environments suited for team projects under deadline.

Key behaviors:

- `uv init` scaffolds the project.
- Locking syncing can be automatic (e.g., `uv run` can ensure environment is synced before execution).

4.2 LangGraph for Agentic Workflow Orchestration

We chose LangGraph because:

- It models the application as a stateful graph with explicit nodes and edges
- It supports common workflow vs. agent patterns and structured routing
- It supports state definitions via Pydantic, improving correctness in multi-agent pipelines

4.3 Pydantic for Schemas & Validation

Pydantic ensures:

- strict typing and validation for:
- search results
- evidence bundles
- outputs shown in the UI
- fewer hidden runtime failures (critical for demos and grading)

4.4 PyTorch + Hugging Face + PEFT/LoRA for Fine-Tuning

We fine-tune a base model with LoRA to:

- efficiently adapt behavior with limited compute
- produce a Writer model that prefers “I don’t know” over hallucination (to be measured in evaluation)

4.5 Streamlit for UI

Streamlit gives:

- fast iteration for a polished demo
- simple UX for showing:
- chat transcript
- “evidence panel”
- run metadata (latency, caching hits, #sources)

5. Repository Structure (Standard, Complete)

```
agentic-historian/
└── README.md
└── LICENSE
└── pyproject.toml
└── uv.lock
└── .python-version
└── .gitignore
└── .env.example
└── ruff.toml
└── mypy.ini
└── pytest.ini
└── docs/
    └── technical_design.md
    └── evaluation_plan.md
    └── demo_script.md
└── data/
    └── raw/
    └── processed/
    └── README.md
└── models/
    └── base/
    └── lora_adapters/
└── notebooks/
    └── 01_truthfulqa_eda.ipynb
    └── 02_lora_sft_experiments.ipynb
    └── 03_eval_analysis.ipynb
└── scripts/
    └── run_app.sh
    └── run_graph_cli.sh
    └── train_lora.sh
    └── eval.sh
└── src/
    └── agentic_historian/
        ├── __init__.py
        ├── config.py
        ├── app/
        │   └── streamlit_app.py
        ├── graph/
        │   ├── state.py
        │   ├── graph.py
        │   └── nodes/
        │       ├── router.py
        │       ├── researcher.py
        │       ├── fact_analyst.py
        │       └── writer.py
        ├── tools/
        │   └── google_search.py
        └── llm/
            ├── prompts.py
            ├── loader.py
            └── generation.py
        └── training/
            ├── data_prep.py
            ├── train_lora.py
            └── eval_truthfulqa.py
        └── utils/
            ├── logging.py
            ├── cache.py
            └── text.py
        └── types/
            └── schemas.py
└── tests/
    └── test_router.py
    └── test_google_tool.py
    └── test_fact_analyst.py
    └── test_writer_constraints.py
    └── test_graph_smoke.py
```

```
■■■ .github/
    ■■■ workflows/
        ■■■ ci.yml
```

6. MLOps & Evaluation Plan (Minimum Viable, Grade-Friendly)

6.1 Offline Evaluation (Model Behavior)

- Use TruthfulQA validation split to evaluate:
- truthfulness-aligned answer preference
- refusal correctness for unanswerable questions
- Metrics:
 - exact/substring match for known answers (when applicable)
 - “refusal rate” on unanswerable prompts
 - qualitative rubric: “evidence grounded / not grounded”

6.2 Online Evaluation (End-to-End App)

- Create 20–30 historical fact queries:
- famous events/dates
- contested claims
- ambiguous questions
- Score:
 - evidence quality
 - citation correctness (IDs map to displayed sources)
 - uncertainty handling in contested/insufficient cases

7. Risk Register & Mitigations

1. No-scraping limitation reduces depth

- Mitigation: explicit “Evidence Limitation” note; rely on multiple sources; domain-restricted search.

2. Search results may conflict

- Mitigation: Fact Analyst labels “contested”; Writer must present both sides with citations.

3. Hallucinations despite fine-tuning

- Mitigation: strict Writer prompt + schema validation + “must cite evidence IDs” check.

4. Latency / quota limits

- Mitigation: caching, query minimization, and demo-safe fallback examples.

8. Run Instructions (Target UX)

- Install dependencies: uv sync
- Run app: uv run streamlit run src/agentic_historian/app/streamlit_app.py
- Train LoRA: uv run python -m agentic_historian.training.train_lora ...
- Evaluate: uv run python -m agentic_historian.training.eval_truthfulqa ...